

Naveen Kumar Rajesh

San Diego, CA • (858)-705-9378 • messinavu10@gmail.com • [linkedin.com/in/messinavu10](https://www.linkedin.com/in/messinavu10) • naveen.work • github.com/Messinavu10

PROFESSIONAL EXPERIENCE

VITAL RECORDS CONTROL (VRC) • San Diego, CA

August 2026 – Present

AI Engineer

- Shaping company-wide AI strategy and roadmap with the CEO and CTO for client products and internal operations.
- Migrating the AskTuring platform to AWS (EKS, ECS, RDS) with RunPod GPU compute, scaling its databases and vector store and integrating with VRC's systems and CI/CD to serve 100K+ customers.
- Enabling agentic access to VRC's internal platforms by building MCP servers that connect existing products to AskTuring, letting employees work across every tool from a single AI interface.

ASKTURING.AI (Acquired by VRC) • San Diego, CA

October 2025 – August 2026

Founding AI Engineer

- Architected the platform's MCP layer, delivering a built-in server and a third-party connector marketplace with prefix-based tool dispatch and OAuth 2.1; expanding agent capabilities and driving a 25% increase in enterprise adoption.
- Led a team to build an agentic text-to-SQL pipeline with an LLM-as-judge retry system, achieving ~85% Spider accuracy and reducing time-to-query readiness ~35%, accelerating customer acquisition.
- Designed and built a hybrid STM/LTM memory system (PostgreSQL rolling-buffer STM, Cohere-embedded LTM in pgvector, LLM-driven consolidation & recency decay), cutting LLM cost 25%.

CHAPMAN UNIVERSITY • Irvine, CA

March 2024 – December 2024

Data Scientist – Contract

- Engineered automated ETL and validation pipelines in Python (Pandas, NumPy) and SQL to process 250M+ prescription records, cutting data preparation time 40% and enabling time-series analysis of opioid trends across California.
- Developed logistic regression and ARIMA models, uncovering 10+ prescribing patterns and forecasts used by Chapman's opioid research team for public-health analysis.

ARMORY INC • San Diego, CA

July 2022 – August 2023

Software Engineer – Contract

- Built Prometheus and Grafana observability across 50+ microservices on AWS EKS, reducing debugging time by 30%.
- Streamlined CI/CD pipelines by containerizing Jenkins and managing infrastructure with Terraform, improving release efficiency by 35% and supporting faster iteration cycles.

EDUCATION

University of California, Irvine

December 2024

Master of Data Science

University of San Diego, Shiley-Marcos School of Engineering

May 2023

B.S. Computer Science

CERTIFICATIONS

[Databricks Data Engineer – Professional](#) • [AWS Machine Learning Engineer – Associate](#) • [GCP Machine Learning Engineer](#)

PROJECTS

Agent Continuous-Learning Platform | Private Repo

- Architected a multi-tenant Go platform that ingests AI agent traces, extracts learnings via 5 strategies, and serves hybrid vector + lexical retrieval with human-in-the-loop review, plus a Python SDK for LangGraph, CrewAI, and Pydantic AI.

Real-Time Media Enrichment Platform | [GitHub](#)

- Built a near-real-time media enrichment pipeline (Kafka, Spark Structured Streaming, Databricks) that catalogs, quality-checks, and AI-enriches text, image, and audio uploads with sentiment, classification, and entity extraction.

FIA Formula 1 Regulations RAG Agent | [GitHub](#)

- Built a LangGraph agent with 5 specialized tools over Pinecone-indexed FIA regulations (2024–2026) for rule search, cross-year comparison, and penalty lookup, with RAGAS evaluation (faithfulness, context precision/recall) in CI.

TECHNICAL SKILLS

Programming Languages & Tools: Python, Go, JavaScript, SQL, R, Bash, Swift, Git, Claude Code, Cursor

LLMs & Agent Systems: Claude, GPT, Gemini, Cohere, LangChain, LangGraph, Langfuse, MCP, RAG, vector embeddings

Backend, Data & ML: FastAPI, Pydantic, Async, DuckDB, PostgreSQL, Redis, Pandas, MLflow, Databricks, A/B testing

Cloud & DevOps: Docker, K8s, AWS (S3, EKS, ECS, RDS, SageMaker), GCP (VertexAI, CloudRun), Terraform, Grafana